

SQL - Basics

Davood Rafiei

Original material Copyright 2001-2022
(some material from textbooks and other instructors)



SQL Overview

- Structured Query Language
 - standard query language for relational system
 - developed in IBM Almaden (system R)
- Some features
 - Declarative: specify the properties that should hold in the result, not how to obtain the result
 - Complex queries have procedural elements
 - Set/Bag semantics
- International Standards
 - SQL1 (1986)
 - SQL2 (SQL-92)
 - SQL3 (SQL-99)
 - Later extensions in 2003, 2006, 2008, 2011, 2016



SQL Overview (cont)

- Data definition component
 - CREATE TABLE table-name (col-defs, constraints)
 - DROP TABLE table-name
 - ALTER TABLE table-name action
 - ✓ modifies the definition of a table where action is
 - ✓ ADD (col-def)
 - ✓ MODIFY (col-def)
 - ✓ ADD constraint
 - ✓ etc.
- Data update component
 - INSERT INTO table-name ...
 - DELETE FROM table-name ...
 - UPDATE table-name ...
- Query Language



Example Tables

branch (bname, address, city, assets)

customer(cname, street, city)

deposit(accno, cname, bname, balance)

loan(accno, cname, bname, amount)

```
CREATE TABLE branch (  
    bname      CHAR(20),  
    address    CHAR(30),  
    city       CHAR(15),  
    asset      INT  
);
```

SQLite command

(in SQLite, a semicolon
always marks the end of a
SQL statement.)



Simple Queries

branch (bname, address, city, assets)

- Find the names of all branches with assets greater than \$2,500,000.

```
SELECT bname  
FROM    branch  
WHERE   assets > 2500000
```

- Find the names of all branches in Edmonton with assets greater than \$2,500,000.

```
SELECT bname  
FROM    branch  
WHERE   assets>2500000 AND city='Edmonton'
```

- Simple predicates can be combined using AND, OR, NOT.



Querying two Relations

customer(cname, street, city)

deposit(accno, cname, bname, balance)

- List the name and the city of every customer who has an account with balance over \$2,000.

```
SELECT customer.cname, city
FROM    customer, deposit
WHERE   balance > 2000
AND     customer.cname = deposit.cname
```



Queries With Tuple Variables

loan(accno, cname, bname, amount)

deposit(accno, cname, bname, balance)

- Find customers who have both loans and deposits.

```
SELECT  loan.cname
FROM    loan, deposit
WHERE   loan.cname = deposit.cname
```

- Equivalently using tuple variables:

```
SELECT  l.cname
FROM    loan l, deposit d
WHERE   l.cname = d.cname
```

- Range variables are really needed if the same relation appears twice in the FROM clause.



A Simple Evaluation Alg.

```
SELECT ...  
FROM   R1 r1, R2 r2, ...  
WHERE  C
```

- *Tuple variables* r_1, r_2, \dots respectively range over rows of R_1, R_2, \dots
- Evaluation strategy:
 - FROM clause produces Cartesian product of listed tables
 - WHERE clause produces table containing only rows satisfying condition
 - SELECT clause retains listed columns

```
for every tuple  $r_1$  in  $R_1, r_2$  in  $R_2, \dots$   
  let  $r := r_1, r_2, \dots$   
  if  $C(r)$  then  
    output the desired columns
```



From SQL to Relational Algebra

```
SELECT l.cname  
FROM   loan l, deposit d  
WHERE  l.cname = d.cname
```

Equivalent to:

$$\begin{aligned} temp &= \rho_{(a1, cn1, bn1, b1, a2, cn2, bn2, a2)} (loan \times deposit) \\ \pi_{cn1} \sigma_{cn1 = cn2} (temp) \end{aligned}$$

This is a simple evaluation algorithm for SELECT.



Queries With Set/Bag Results

- Find all cities of customers.
 - **SELECT** city
FROM customer
- Result?
- To get rid of duplicates, we need
 - **SELECT** **DISTINCT** city
FROM customer



Duplicates

- Duplicate rows not allowed in a relation
- However, duplicate elimination from query result is costly and not automatically done; it must be explicitly requested:

```
SELECT DISTINCT .....  
FROM .....
```



Working with Strings

- Equality and comparison operators apply to strings (based on lexical ordering)
 - E.g. WHERE *cname* < 'P'
- Concatenate operator applies to strings
 - E.g. WHERE *bname* || '--' || *address* =
- Expressions can also be used in SELECT clause
 - E.g. SELECT *bname* || '--' || *address* AS NameAdd
FROM branch



Partial Matching

customer(cname, street, city)

- Find every customer whose address starts with “Computing Science”.

```
SELECT  *  
FROM    customer  
WHERE    address LIKE 'Computing Science%'
```

- Expression: col-name [NOT] LIKE pattern
 - Pattern may include wildcard characters ‘%’ matching any string and ‘_’ (underscore) matching any single character.



Naming the Results

deposit(accno, cname, bname, balance)

- For every deposit holder who has over \$1000, find the customer name and the balance over \$1000.

```
SELECT  cname, (balance - 1000) as bal
FROM    deposit
WHERE    balance > 1000
```



Ordering the Results

branch (bname, address, city, assets)

- Find the names and assets of all branches with assets greater than \$2,500,000 and order the result in ascending order of asset values.

```
SELECT  bname, assets
FROM    branch
WHERE    assets > 2500000
ORDER BY assets
```

- Default is ascending order; a descending order can be specified by the DESC keyword.



Queries Involving Set Operators

- set union : Q1 UNION Q2
 - the set of tuples in Q1 or Q2
- set difference : Q1 EXCEPT Q2
 - the set of tuples in Q1 but not in Q2
- set intersection : Q1 INTERSECT Q2
 - the set of tuples in both Q1 and Q2
- Q1 and Q2 must be union-compatibles
 - same number/types of attributes.



Queries With Set Operators

- List deposit-holders who have no loans.

```
(SELECT  cname
FROM    deposit)
```

EXCEPT

```
(SELECT  cname
FROM    loan)
```

- List cities where there is either a customer or a branch.

```
(SELECT  city
FROM    customer)
```

UNION

```
(SELECT  city
FROM    branch)
```



Queries With Set Operators

- Find all cities that have both customers and branches in.

```
(SELECT city  
FROM      customer)
```

INTERSECT

```
(SELECT city  
FROM      branch)
```

- Find every city that has a branch but no customer.



Queries Over Multiple Relations

branch (bname, address, city, assets)

customer(cname, street, city)

deposit(accno, cname, bname, balance)

```
SELECT branch.bname, assets
FROM    branch, customer, deposit
WHERE   customer.city = 'Jasper'
AND     customer.cname = deposit.cname
AND     deposit.bname = branch.bname
```

- What does the query do?

Find the name and asset of every branch that has a deposit account holders who lives in Jasper.



Queries With Nested Structures

- Queries within the WHERE clause of an **outer** query

SELECT

FROM

WHERE OPERATOR (SELECT ... FROM ... WHERE)

- There can be multiple levels of nesting
- Operators: **IN**, **(NOT) EXISTS**, **< ALL**, ...
- **Avoid nesting as much as possible.**



Nested Structures Using “IN”

```
SELECT  
FROM  
WHERE expr | (expr list) IN  
                                (set of values)
```

- E.g.
... **WHERE** *province* **IN** ('AB', 'BC')
... **WHERE** *province* **NOT IN** ('AB', 'BC')



Example

deposit(accno, cname, bname, balance)

```
SELECT cname
FROM    deposit
WHERE    bname IN
           (SELECT    bname
            FROM    deposit
            WHERE    cname = 'John Doe' )
```

- What does the query do?

Find every customer who has a deposit in some branch at which John Doe has a deposit.



The Same Example Without Nesting

deposit(accno, cname, bname, balance)

```
SELECT    d1.cname
FROM      deposit d1, deposit d2
WHERE     d2.cname = 'John Doe'
AND       d1.bname = d2.bname
```

NOTE: avoid nesting as much as possible.



Nested Structures Using “< ALL”, ...

SELECT
FROM

WHERE *expr* < **ALL** (set of values)

- Other forms: “<= ALL”, “= ALL”, “>= ALL”, “> ALL”
- Op ALL (set of values) evaluates to **true** iff the comparison evaluates to true for every value in the set.
- Op SOME (set of values) evaluates to **true** iff the the comparison evaluates to true for at least one value in the set.



Nested Structures Using “>ALL”, ...

branch (bname, address, city, assets)

- Find branches that have assets greater than the assets of all branches in Calgary.

```
SELECT    bname
FROM      branch
WHERE      assets > ALL
              (SELECT assets
               FROM    branch
               WHERE   city = 'Calgary')
```



Nested Structures Using “EXISTS”

SELECT

FROM

WHERE (NOT) EXISTS (SELECT ...)

- **EXISTS (SELECT ...)**
evaluates to true iff the result of the subquery contains at least one row.
- The expression is evaluated for every iteration of the outer query.



“EXISTS” Construct Example

branch (bname, address, city, assets)

customer(cname, street, city)

- Find the names of customers who live in a city with no bank branches.


```
SELECT  cname
FROM    customer
WHERE NOT EXISTS (SELECT *
                  FROM    branch
                  WHERE    customer.city=branch.city)
```

- Find the names of customers who live in a city which has a bank branch.
 - Change NOT EXISTS to EXISTS.
(can also write it using join)



Division

- *Query type*: Find the subset of items in one set that are related to *all* items in another set
- *Example*: Find customers who have deposit accounts in *all* branches.



$$\pi_{cname, bname}(\text{deposit}) / \pi_{bname}(\text{branch})$$



Division

- *Strategy for implementing division in SQL:*
 - Find set, A, of all branches in which a particular customer, c, has a deposit account.
 - Find set, B, of all branches.
 - Output c if $A \supseteq B$, or, equivalently, if $B-A$ is empty



Division – SQL Solution

branch (bname, address, city, assets)

deposit(accno, cname, bname, balance)

```
SELECT c.cname
FROM customer c
WHERE NOT EXISTS
    (SELECT b.bname                -- set B of all branches
     FROM branch b
     EXCEPT
     SELECT d.bname                -- set A of branches in which
                                     -- customer c has a deposit account
     FROM deposit d
     WHERE d.cname=c.cname        -- global variable
    )
```



Division – 2nd SQL Solution

- Find customers who have deposit accounts in all branches.
 - Same as “find all customers such that there is no branch where they do not have deposits in.”

```
SELECT c.cname
FROM   customer c
WHERE  NOT EXISTS
```

```
    (SELECT  bname
     FROM    branch b
     WHERE   NOT EXISTS
```

```
        (SELECT *
         FROM   deposit
         WHERE  b.bname=deposit.bname
         AND    c.cname = deposit.cname))
```

branches where
c has no deposits in

Deposits of c in
branch b



Division – Another Example

- Find professors who have taught courses in *all* departments.
- *Strategy for implementing division in SQL:*
 - Find set, A, of all departments in which a particular professor, p , has taught a course
 - Find set, B, of all departments
 - Output p if $A \supseteq B$, or, equivalently, if $B - A$ is empty



Division – SQL Solution

Professor (*Id, Name, DeptId*)

Department (*DeptId, Name*)

Course (*DeptId, CrsCode, CrsName, Descr*)

Teaching (*ProfId, CrsCode, Semester*)

```
SELECT P.Id
FROM Professor P
WHERE NOT EXISTS
    (SELECT D.DeptId                -- set B of all dept Ids
      FROM Department D
      EXCEPT
      SELECT C.DeptId                -- set A of dept Ids of depts in
                                     -- which P has taught a course
      FROM Teaching T, Course C
      WHERE T.ProfId=P.Id          -- global variable
            AND T.CrsCode=C.CrsCode)
```



SQL and SQLite

- SQLite: pretty modern but light language
 - Released in 2000 (Oracle was 1977, DB2 was 1983)
 - It is light (intended for small devices)
 - Tightly integrated with applications (software library rather than a stand-alone system/process to communicate)
 - Locks the whole file/database during writes (not a good choice for write-intensive/concurrent transactions)
 - Supports both in-memory and on-disk databases
- Differences with SQL-92
 - Does not support op ALL in nested queries
 - Does not support RIGHT OUTER JOIN and FULL OUTER JOIN
 - Views are read-only
 - ALTER TABLE command is very limited
 - Foreign keys constraints are not enforced by default
 - More left to be explored!



Basic Data Types in SQLite

- **INTEGER** 1, 2, 3, 4, 6, 8 bytes integer
(depending on the magnitude)
- **REAL** 8 bytes floating point number
- **TEXT** stored using database encoding (e.g. UTF-8)
- **BLOB** stored as is
- **NULL** null value
- **NUMERIC** can store all other types; stored value is converted to
INTEGER or REAL if the conversion is lossless and reversible



Other Types in SQLite

- | | |
|-----------------------|---------|
| ■ INT, SMALLINT | INTEGER |
| ■ CHAR(n), VARCHAR(n) | TEXT |
| ■ DOUBLE, FLOAT | REAL |
| ■ DECIMAL | NUMERIC |
| ■ DATE, DATETIME | NUMERIC |

